

5.

collaborator:

B11901040 項蓮均

$$G(x) = \text{sign}\left(\sum_{t=1}^{2M+1} g_t(x)\right)$$

$$E_{\text{out}}(G) = \mathbb{I}[y \neq \text{sign}\left(\sum_{t=1}^{2M+1} g_t(x)\right)]$$

at least $M+1$ classifiers $y \neq \text{sign}(g_t(x))$ s.t. $E(G) = 1$

$$\Rightarrow \sum_{t=1}^{2M+1} e_t \geq M+1$$

$$\Rightarrow 1 \leq \frac{1}{M+1} \sum_{t=1}^{2M+1} e_t$$

$$\Rightarrow \underline{\underline{E(G) \leq \frac{1}{M+1} \sum_{t=1}^{2M+1} e_t}}$$

6.

collaborator:

B11901040 項蓮均

$$\text{incorrect: } u_n^{(t+1)} \leftarrow u_n^{(t)} * \sqrt{\frac{1-E_t}{E_t}}$$

$$\text{correct: } u_n^{(t+1)} \leftarrow u_n^{(t)} / \sqrt{\frac{1-E_t}{E_t}}$$

$$U_{t+1} = \sum_{n=1}^N u_n^{(t+1)}$$

$$= \sum_{\text{incorrect } n} \left(u_n^{(t)} * \sqrt{\frac{1-E_t}{E_t}} + \sum_{\text{correct } n} u_n^{(t)} / \sqrt{\frac{1-E_t}{E_t}} \right)$$

$$= \sqrt{\frac{1-E_t}{E_t}} \sum_{\text{incorrect } n} u_n^{(t)} + \sqrt{\frac{E_t}{1-E_t}} \sum_{\text{correct } n} u_n^{(t)}$$

$$= \sqrt{\frac{1-E_t}{E_t}} U_t * E_t + \sqrt{\frac{E_t}{1-E_t}} U_t (1-E_t)$$

$$= U_t \left(\sqrt{E_t(1-E_t)} + \sqrt{E_t(1-E_t)} \right)$$

$$= 2\sqrt{E_t(1-E_t)} U_t$$

$$\therefore \frac{U_{t+1}}{U_t} = 2\sqrt{E_t(1-E_t)} \quad \text{Q. E. D.}$$

7.

no collaborators

$$\min_{\alpha} \frac{1}{N} \sum_{n=1}^N ((y_n - s_n) - \alpha g_t(\underline{x}_n))^2$$

$$\frac{1}{N} \sum_{n=1}^N ((y_n - s_n) - \alpha g_t(\underline{x}_n))^2$$

$$= \frac{1}{N} \sum_{n=1}^N ((y_n - s_n)^2 - 2\alpha (y_n - s_n) g_t(\underline{x}_n) + \alpha^2 g_t^2(\underline{x}_n))$$

$$\frac{\partial}{\partial \alpha} \frac{1}{N} \sum_{n=1}^N ((y_n - s_n)^2 - 2\alpha (y_n - s_n) g_t(\underline{x}_n) + \alpha^2 g_t^2(\underline{x}_n)) = 0$$

$$\Rightarrow \sum_{n=1}^N (-2(y_n - s_n) g_t(\underline{x}_n) + 2\alpha g_t^2(\underline{x}_n)) = 0$$

$$\Rightarrow 2\alpha \sum_{n=1}^N g_t^2(\underline{x}_n) = 2 \sum_{n=1}^N (y_n - s_n) g_t(\underline{x}_n)$$

$$\Rightarrow \alpha = \frac{\sum_{n=1}^N (y_n - s_n) g_t(\underline{x}_n)}{\sum_{n=1}^N g_t^2(\underline{x}_n)}$$

From the solution of linear regression, $((y_n - s_n) - g_t(\underline{x}_n))$ are orthogonal to $g_t(\underline{x}_n)$

$$\Rightarrow \sum_{n=1}^N ((y_n - s_n) - g_t(\underline{x}_n)) \cdot g_t(\underline{x}_n) = 0$$

$$\Rightarrow \sum_{n=1}^N (y_n - s_n) g_t(\underline{x}_n) - \sum_{n=1}^N g_t^2(\underline{x}_n) = 0$$

$$\Rightarrow \sum_{n=1}^N (y_n - s_n) g_t(\underline{x}_n) = \sum_{n=1}^N g_t^2(\underline{x}_n)$$

$$\therefore \alpha_1 = \frac{\sum_{n=1}^N (y_n - s_n) g_t(\underline{x}_n)}{\sum_{n=1}^N g_t^2(\underline{x}_n)} = 1 \quad \text{Q.E.D.}$$

8.

perform linear regression on $\{g_t(\underline{x}_n), y_n - s_n\}_{n=1}^N$

collaborator:

B11901040 項蓮均

$$\min_{\eta} \frac{1}{N} \sum_{n=1}^N ((y_n - s_n) - \eta g_t(\underline{x}_n))^2$$

$$\stackrel{\frac{\partial}{\partial \eta}}{\Rightarrow} \frac{1}{N} \sum_{n=1}^N 2(-g_t(\underline{x}_n))((y_n - s_n) - \eta g_t(\underline{x}_n)) = 0$$

$$\Rightarrow \sum_{n=1}^N g_t(\underline{x}_n)((y_n - s_n) - \eta g_t(\underline{x}_n)) = 0$$

$$\eta = \alpha_t$$

$$\text{update } s_n \leftarrow s_n + \alpha_t g_t(\underline{x}_n)$$

$$\therefore \sum_{n=1}^N g_t(\underline{x}_n)(y_n - (s_n + \alpha_t g_t(\underline{x}_n))) = 0$$

$$\Rightarrow \sum_{n=1}^N g_t(\underline{x}_n)(y_n - s_n) = 0 \quad \text{Q.E.D.}$$

9.

collaborator:

B11201009 黃勤元

\therefore one hidden layer including the output neuron

$\therefore L = 2$

$$\begin{cases} s_j^{(2)} = \sum_{i=0}^{d^{(1)}} w_{ij}^{(2)} x_i^{(1)} \\ s_j^{(1)} = \sum_{i=0}^{d^{(0)}} w_{ij}^{(1)} x_i^{(0)}, \quad s_{j+1}^{(1)} = \sum_{i=0}^{d^{(0)}} w_{i,j+1}^{(1)} x_i^{(0)} \\ \delta_j^{(2)} = -2(y_n - s_j^{(2)}) \end{cases}$$

$$w_{ij}^{(0)} = 0.5 \quad \forall i \text{ and } 1 \leq j < d^{(1)}$$

$$\begin{aligned} \delta_j^{(1)} &= \sum_k \delta_k^{(2)} w_{jk}^{(2)} \tanh'(s_j^{(1)}) \\ &= \sum_k \delta_k^{(2)} w_{j+1,k}^{(2)} \tanh'(s_{j+1}^{(1)}) \\ &= \delta_{j+1}^{(1)} \end{aligned}$$

for the backprop algorithm,

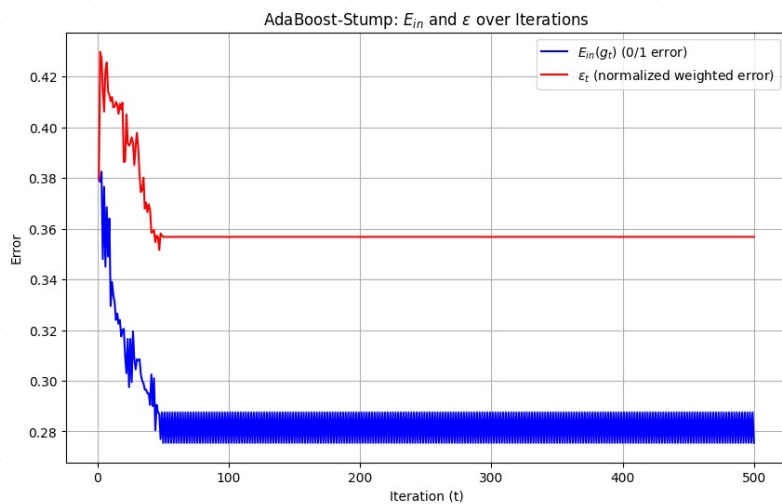
$$w_{ij}^{(2)} = w_{ij}^{(2)} - \eta x_i^{(1)} \delta_j^{(2)}$$

$$\begin{aligned} w_{ij}^{(1)} &= w_{ij}^{(1)} - \eta x_i^{(0)} \delta_j^{(1)} \\ &= w_{i,j+1}^{(1)} - \eta x_i^{(0)} \delta_{j+1}^{(1)} \end{aligned}$$

$$= w_{i,j+1}^{(1)}$$

Q.E.D.

10. collaborators: B11901073 林禹融



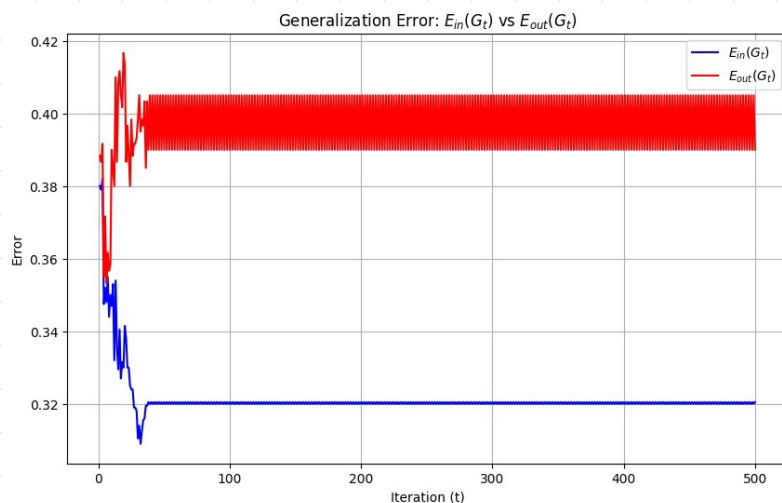
The AdaBoost algorithm reduces $E_{in}(g_t)$ and ϵ_t quickly and remain same value respectively.

```

9 class DecisionStump:
10     def __init__(self):
11         self.feature_index = None
12         self.threshold = None
13         self.sign = None
14
15     def train(self, X, y, weights):
16         _, n = X.shape
17         best_error = float('inf')
18
19         # Search for the best dimension, threshold, and sign
20         for feature_index in range(n):
21             thresholds = np.unique(X[:, feature_index])
22             for threshold in thresholds:
23                 for sign in [-1, 1]:
24                     predictions = sign * np.sign(X[:, feature_index] - threshold)
25                     error = np.sum(weights * (predictions != y))
26                     if error < best_error:
27                         best_error = error
28                         self.feature_index = feature_index
29                         self.threshold = threshold
30                         self.sign = sign
31
32         return best_error
33
34     def predict(self, X):
35         predictions = self.sign * np.sign(X[:, self.feature_index] - self.threshold)
36         return predictions
37
38 # AdaBoost Algorithm
39 def adaBoost(X_train, y_train, X_test, y_test, T=500):
40     m = X_train.shape[0]
41     weights = np.ones(m) / m
42     classifiers = []
43     alphas = []
44
45     # Metrics for plotting
46     E_in = []
47     epsilons = []
48
49     for t in range(T):
50         stump = DecisionStump()
51         error = stump.train(X_train, y_train, weights)
52
53         # Avoid divide-by-zero
54         if error == 0:
55             error = 1e-10
56
57         # Compute alpha
58         alpha = 0.5 * np.log((1 - error) / error)
59         alphas.append(alpha)
60         classifiers.append(stump)
61
62         # Update weights
63         predictions = stump.predict(X_train)
64         weights *= np.exp(-alpha * y_train * predictions)
65
66         # Compute 0/1 error (E_in) and epsilon
67         E_in.append(zero_one_loss(y_train, np.sign(sum(alpha * clf.predict(X_train) for clf, alpha in zip(classifiers, alphas))))
68         epsilons.append(error)
69
70         print(f"Iteration {t+1}: Error = {error:.4f}, Alpha = {alpha:.4f}, E_in = {E_in[-1]:.4f}")
71
72     # Final testing error
73     final_predictions = np.sign(sum(alpha * clf.predict(X_test) for clf, alpha in zip(classifiers, alphas)))
74     E_out = zero_one_loss(y_test, final_predictions)
75
76     print(f"Final Test Error (E_out): {E_out:.4f}")
77
78     return E_in, epsilons
79

```

11. collaborators: B11901073 林禹融

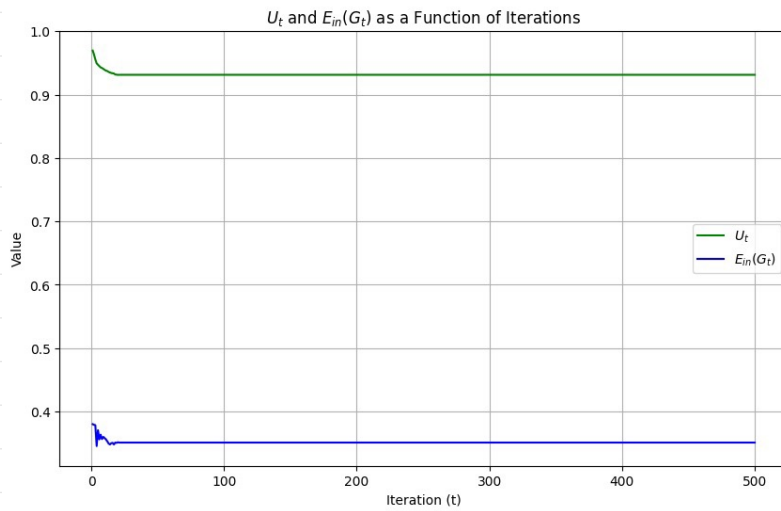


① Although $E_{in}(G_t)$ is relatively high initially, it decreases fast and remains around 0.32 after about 50 iterations.

② $E_{out}(G_t)$ oscillates between about 0.39 and 0.41 after about 50 iterations. Each new weak classifier added to G_t contributes a small correction based on its d_t . If the weak classifier's corrections are small and do not generalize well to the test data, they may slightly shift the decision boundary of G_t in an inconsistent manner. However, the mean value of the oscillation remains about 0.395 after about 50 iterations.

① & ② indicate that after a certain number of iterations, the model reaches its generalization capacity.

12. no collaborators



It decreases steeply in the first few iterations and remains at about 0.93, which suggest that AdaBoost algorithm quickly adjusts the sample weights and converges efficiently. However, it doesn't reach perfect accuracy.